

Week 14, Lec 27

# Transaction and Concurrency Control

Part 1

# Contents

- **Transaction control**
- **Data Concurrency and Consistency in a Multiuser Environment**
- **Locking**

# ***A Banking Transaction***

## **Transaction Begins**

```
UPDATE savings_accounts  
  SET balance = balance - 500  
  WHERE account = 3209;
```

Decrement Savings Account

```
UPDATE checking_accounts  
  SET balance = balance + 500  
  WHERE account = 3208;
```

Increment Checking Account

```
INSERT INTO journal VALUES  
  (journal_seq.NEXTVAL, '1B'  
   3209, 3208, 500);
```

Record in Transaction Journal

```
COMMIT WORK;
```

End Transaction

## **Transaction Ends**

# Transactions - Rationale

- Consider two clients booking airline tickets
- There are 2 seats left on a flight
- Client A wants 2 seats:
  - time 12:02 makes initial request
  - 12:06 confirms purchase through booking form
  - 12:08 authorises credit card payment
- Client B wants 2 seats:
  - time 12:03 makes initial request
  - 12:05 confirms purchase through booking form
  - 12:09 authorises credit card payment
- Situation needs careful control

# Some Possibilities

- Clients A and B are both told 2 seats are free in initial enquiries
- B confirms purchase before A
  - But A may still proceed
- A attempts credit card debit first
  - If successful A secures tickets at 12:08
- B then attempts credit card debit
  - If successful B secures tickets at 12:09
    - potentially over-writing A's tickets
    - A has paid for tickets no longer his/hers

# Requirements 1

- When client A beats B in the initial enquiry:
  - they should form a queue (serialisability)
  - B must wait for A to finish
- Different kinds of finish for A:
  - successful
    - completes booking form
    - makes credit card debit
    - store results (commit)
      - number of seats available is now zero
    - write transaction log and finish
    - B cannot proceed with purchase as no tickets left

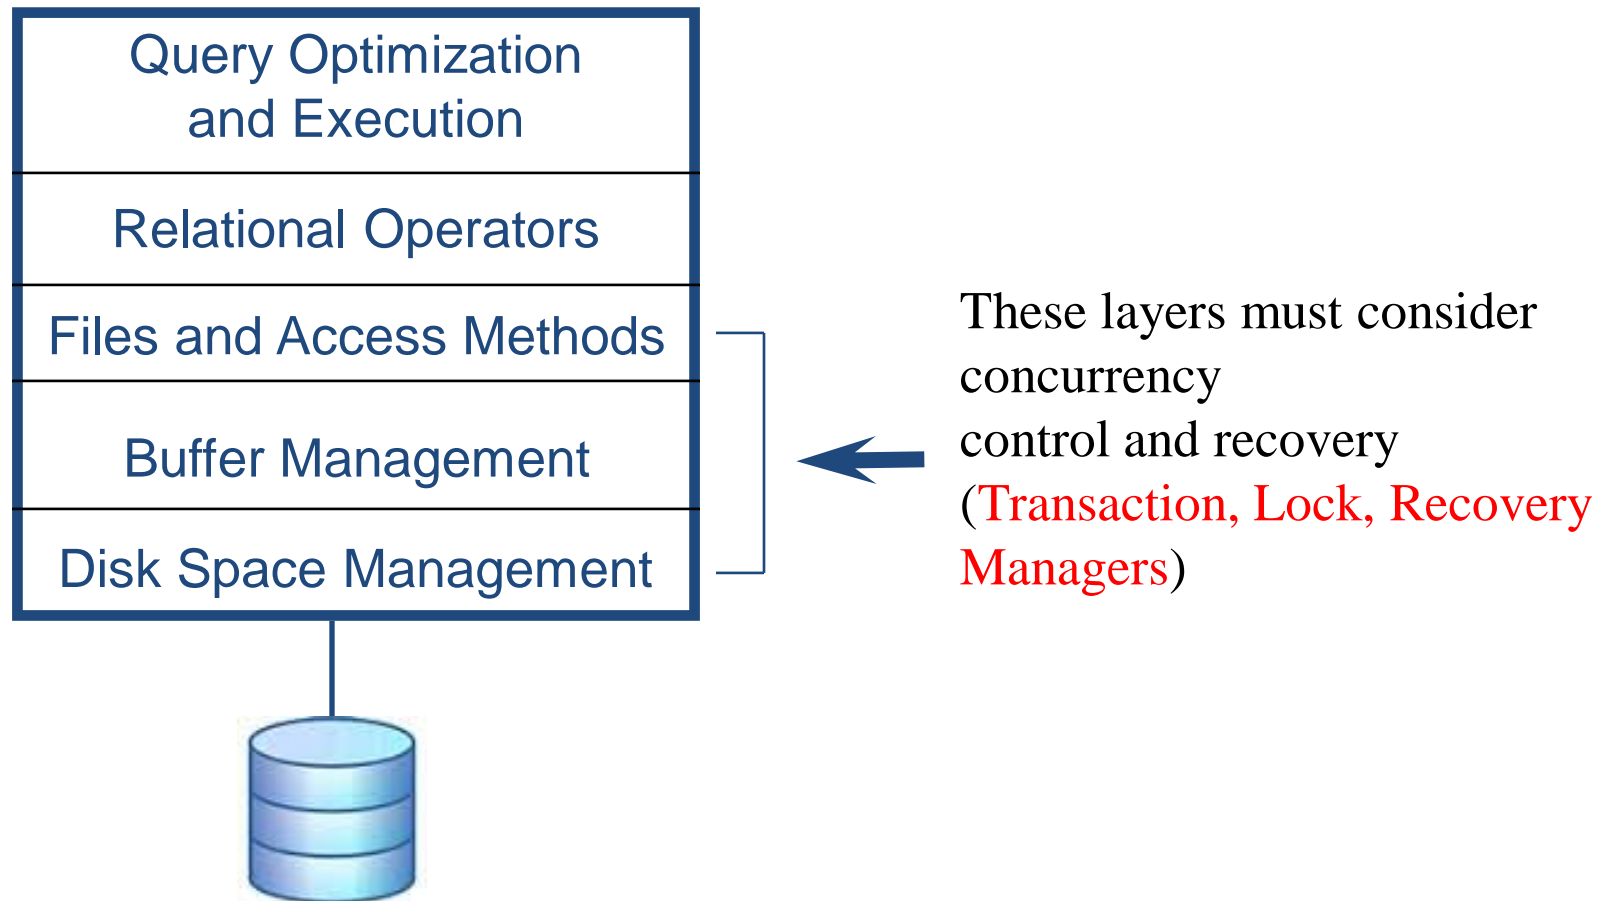
# Requirements 2

- unsuccessful

- may not complete booking form
- may not have funds on credit card
- undo any database changes (rollback) and finish
  - number of seats available is still 2
- B can now proceed to attempt to purchase the 2 tickets left

- Techniques required to emulate business practice

# Structure of a DBMS





# Transactions and Concurrent Execution

- Transaction - DBMS's abstract view of a user program (or activity):
  - A sequence of **reads** and **writes** of database objects.
  - Unit of work that must **commit** or **abort** as an **atomic unit**
- Transaction Manager controls the execution of transactions.
- User's program logic is invisible to DBMS!
  - Arbitrary computation possible on data fetched from the DB
  - The DBMS only sees data read/written from/to the DB.
- Challenge: provide atomic transactions to concurrent users!
  - Given only the read/write interface.

# Concurrency: Why bother?

- The *latency* argument
  - Latency
    - Average response time
    - Average time taken to complete a transaction
- The *throughput* argument
  - System throughput:
    - Number of transactions executed per time unit
- Both are critical!

# Example of a Fund Transfer

- **Isolation requirement** — if between steps 3 and 6, another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be).

**T1**

1. **read**( $A$ )
2.  $A := A - 50$
3. **write**( $A$ )
4. **read**( $B$ )
5.  $B := B + 50$
6. **write**( $B$ )

**T2**

read( $A$ ), read( $B$ ), print( $A+B$ )

- Isolation can be **ensured trivially** by running transactions **serially**
  - that is, one after the other.
- However, executing multiple transactions concurrently has significant benefits, as we will see later.

# ACID Properties

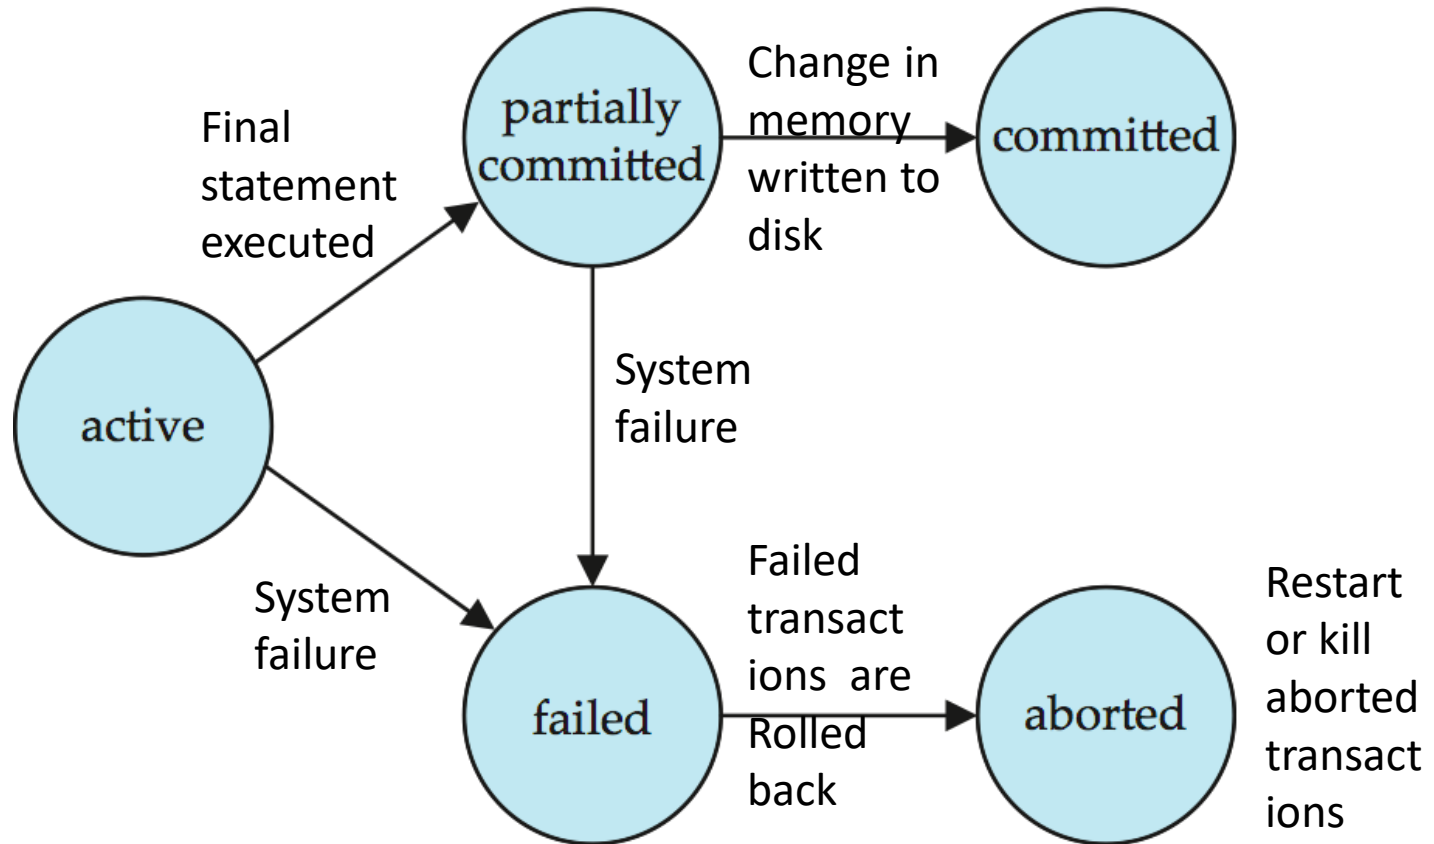
A **transaction** is a **unit of program execution** that accesses and possibly updates various data items. To preserve the **integrity of data** the database system must ensure:

- **Atomicity.** **Either all** operations of the transaction are **properly** reflected in the database **or none** are.
- **Consistency.** **Execution** of a transaction in **isolation preserves** the **consistency** of the database.
- **Isolation.** Although multiple transactions may execute concurrently, **each transaction must be unaware of other concurrently executing** transactions. **Intermediate transaction results must be hidden** from other concurrently executed transactions.
  - That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.
- **Durability.** After a **transaction completes successfully**, the **changes** it has made to the database **persist**, even if there are system failures.

# Transaction State

- **Active** – the **initial state**; the transaction stays in this state while it is executing
- **Partially committed** – after the **final statement** has been **executed**.
- **Failed** – after the discovery that **normal execution** can **no longer proceed**.
- **Aborted** – after the transaction has been **rolled back** and the **database restored to its state prior** to the start of the transaction. Two options after it has been aborted:
  - **restart** the transaction
    - can be done only if no internal logical error
  - **kill** the transaction
- **Committed** – after **successful completion**.

# Transaction State (Cont.)



# Atomicity and Durability

- A transaction ends in one of two ways:
  - *commit* after completing all its actions
    - “commit” is a contract with the caller of the DB
  - *abort* (or *be aborted* by the DBMS) after executing some actions.
    - Or *system crash* while the xact is in progress; treat as abort.
- Two important properties for a transaction:
  - *Atomicity* : Either execute all its actions, or none of them
  - *Durability* : The *effects of a committed transaction* must survive failures.
- DBMS *ensures* the above *by logging* all actions:
  - *Undo* the actions of aborted/failed transactions.
  - *Redo* actions of *committed transactions not yet propagated to disk* when system crashes.

# SQL Transaction commands

- DBMS does **not have an built-in way of knowing** which commands are **grouped** to form a single logical transaction.
- **Some commands**, e.g. **COMMIT** and **ROLLBACK** can **provide boundaries** of transaction.
- **Commit**
  - saves current database state
  - releases resources, locks & savepoints held
  - equivalent to **Save and Exit** in MS Word
- **Rollback**
  - returns database state to that at start of transaction
  - releases resources, locks & savepoints held
  - equivalent to **dismiss/ do not save changes** in MS Word
- **By default**, every SQL statement also **commits implicitly** if it executes **successfully**
  - Implicit commit can be turned off
    - E.g. in SQLPLUS, set autocommit off



# Transactions in SQL

- A transaction is a **logical unit of work** on a database.
- A **group of related operations** that
  - typically comprises a collection of individual actions
    - e.g. in SQL INSERT, UPDATE, DELETE, SELECT
  - must be performed successfully
    - before any changes to the database are finalised.
- Variable size:
  - entire run on SQL\*Plus
    - e.g. spend 2 hours inserting data
  - single command in SQL\*Plus
    - e.g. one insert command
  - one execution of a procedure
    - e.g. one run of add\_patient

# Database Transaction

**A database transaction consists of one of the following:**

- **DML statements which constitute one consistent change to the data**
- **One DDL statement**
- **One DCL statement**

# Oracle Transaction Types

Type	Description
Data manipulation language (DML)	Consists of any number of DML statements that the Oracle server treats as a single entity or a logical unit of work
Data definition language (DDL)	Consists of <b>only one</b> DDL statement
Data control language (DCL)	Consists of <b>only one</b> DCL statement (GRANT, REVOKE)

# Transaction boundaries

A transaction **begins with** the first executable SQL statement.

A transaction **ends with** one of the following events:

- A COMMIT or ROLLBACK statement is issued
- A DDL or DCL statement executes (automatic commit)
- The user exits *iSQL\*Plus*
- The system crashes

# Advantages of COMMIT and ROLLBACK

With COMMIT and ROLLBACK statements, you can:

- Ensure data consistency
- Preview data changes before making changes permanent
- Group logically related operations

# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system. Advantages are:
  - **increased processor and disk utilization**, leading to better transaction *throughput*
    - ▶ E.g. one transaction can be using the CPU while another is reading from or writing to the disk
  - **reduced average response time** for transactions: short transactions need not wait behind long ones.
- **Concurrency control schemes** – mechanisms to **achieve isolation**
  - **Control the interaction among the concurrent transactions** in order to prevent them from destroying the consistency of the database

# Schedules

- **Schedule** – a sequences of instructions that specify the **chronological order** in which instructions of concurrent transactions are executed
  - a schedule for a **set of transactions** must consist of **all instructions** of those transactions
  - must **preserve the order** in which the **instructions appear in each** individual transaction.
- A **transaction that successfully** completes its execution will have a **commit instructions** as the last statement
  - by default transaction assumed to execute commit instruction as its last step
- A transaction **that fails to successfully complete** its execution will have an **abort instruction** as the last statement

# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- Goal:  $A + B$  is “preserved”.
- A **serial schedule** in which  $T_1$  is followed by  $T_2$  :

$A = 100$

$B = 10$

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit



# Schedule 2

- A **serial schedule** where  $T_2$  is followed by  $T_1$

$A = 100$

**B = 10**

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

# Schedule 3

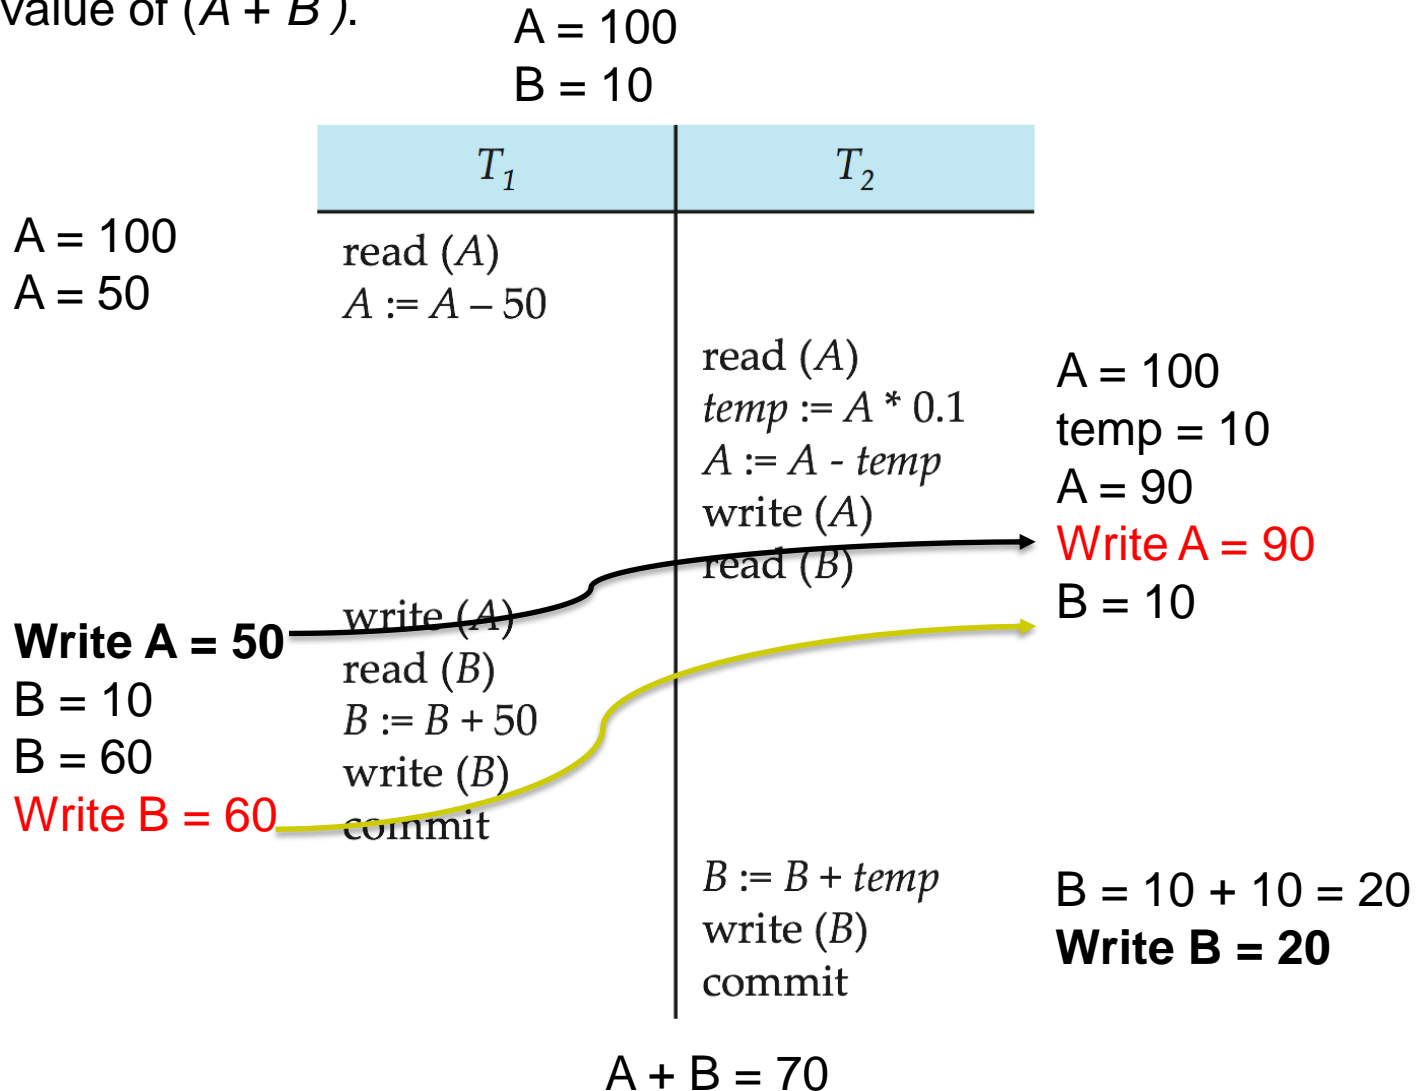
- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is **not a serial schedule**, but it is **equivalent** to **Schedule 1**.

	$A = 100$ $B = 10$	
	$T_1$	$T_2$
$A = 100$ $A = 50$ <b>Write A = 50</b>	$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$	
		$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$
		$A = 50$ $\text{temp} = 5$ $A = 45$ <b>Write A = 45</b>
$B = 10$ $B = 60$ <b>Write B = 60</b>	$\text{read}(B)$ $B := B + 50$ $\text{write}(B)$ $\text{commit}$	
		$\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$ $\text{commit}$
		$B = 60$ $B = 65$ <b>Write B = 65</b>
	$A + B = 110$	

In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.

# Schedule 4

- The following **concurrent** schedule **does not preserve** the value of  $(A + B)$ .



# Serializability

- **Basic Assumption** – Each transaction preserves database consistency.
- Thus serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
  1. **conflict serializability**
  2. **view serializability**

# ***Simplified view of transactions***

- We ignore operations other than **read** and **write** instructions
- We assume that transactions may perform **arbitrary computations** on data **in local buffers in between** reads and writes.
- Our simplified schedules **consist of only read and write** instructions.

# Conflicting Instructions

- Instructions  $I_i$  and  $I_j$  of 2 transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists **some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and at least one of these instructions wrote  $Q$ .**
  1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict.
  2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . **They conflict.**
  3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . **They conflict**
  4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . **They conflict**
- Intuitively, a **conflict** between  $I_i$  and  $I_j$  **forces a (logical) temporal order** between them.
  - If  $I_i$  and  $I_j$  are **consecutive in a schedule** and they **do not conflict**, their results would remain the **same** even if they had been **interchanged** in the schedule.

# Conflict Serializability

- If a schedule  $S$  can be transformed into another schedule  $S'$  by a series of swaps of *non-conflicting instructions*, we say that  $S$  and  $S'$  are **conflict equivalent**.
- We say that a schedule  $S$  is **conflict serializable** if it is **conflict equivalent** to a **serial** schedule

# Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions. Therefore **Schedule 3 is conflict serializable.**

**Transforming details were discussed on board in class. (to be continued)**

$T_1$	$T_2$
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

Schedule 3

$T_1$	$T_2$
read (A) write (A) read (B) write (B)	read (A) write (A) read (B) write (B)

Schedule 6