

Week 14, Lec 27

Transaction and Concurrency Control

Part 2

Contents

- **Transaction control**
- **Data Concurrency and Consistency in a Multiuser Environment**
- **Locking**

Conflict Serializability

- If a schedule S can be transformed into another schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent.
- We say that a schedule S is **conflict serializable** if it is **conflict equivalent** to a **serial** schedule

Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1 , by series of swaps of non-conflicting instructions. Therefore **Schedule 3 is conflict serializable**.

Transforming details were discussed on board in class.

T_1	T_2		T_1	T_2
read (A) write (A)			read (A) write (A)	
	read (A) write (A)	→	read (B) write (B)	
read (B) write (B)				read (A) write (A)
	read (B) write (B)			read (B) write (B)
Schedule 3			Schedule 6	

Conflict Serializability (Cont.)

- Example of a schedule that is **not conflict serializable**:

T_3	T_4
read (Q)	
write (Q)	write (Q)


- We are **unable to swap instructions** in the above schedule **to obtain either** the **serial schedule** $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

Testing for Serializability


- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence graph** — a **directed** graph where the **vertices** are the **transactions (names)**.
- We draw an **arc from T_i to T_j** if the **two transaction conflict, and T_i accessed** the **data item** on which the **conflict** arose **earlier**.
- We **may label the arc by** the **item** that was **accessed**.
- **Edge: $T_i \rightarrow T_j$ means one of 3 conditions holds**
 - T_i runs read(Q) before T_j runs write(Q)
 - T_i runs write(Q) before T_j runs read(Q)
 - T_i runs write(Q) before T_j runs write(Q)

Testing for Serializability


- Example 1 –
Precedence
Graph for
Schedule 1



- Example 2 –
Precedence
Graph for
Schedule 2



- Example 3 –
Precedence
Graph for
Schedule 4




Schedule 1 and Its Precedence Graph

- A **serial** schedule in which T_1 is followed by T_2 :

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Precedence Graph for Schedule 1




Schedule 2 and its Precedence Graph

- A **serial** schedule where T_2 is followed by T_1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Precedence Graph for Schedule 2




Schedule 4 and its Precedence Graph

- The following **concurrent schedule *does not preserve*** the value of $(A + B)$.




T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	 read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Precedence Graph for Schedule 4






Test for Conflict Serializability

- A schedule is **conflict serializable if and only if its precedence graph is acyclic**.
- **Cycle-detection algorithms** exist which take **order n^2 time**, where n is the number of vertices in the graph.
- If precedence graph is **acyclic**, the **serializability order** can be obtained by a **topological sorting** of the graph.
 - This is a **linear order** consistent with the partial order of the graph.
 - For example, a **serializability order** for Schedule A would be
 $T_i \rightarrow T_j \rightarrow T_k \rightarrow T_m$



Test for Conflict Serializability

- *topological sorting* of a directed acyclic graph (DAG)
 - Ordering of vertices such that **if (u,v) is an edge**, then **u appears before v** in the **resulting order**
 - One topological sorting of Schedule (a) would be in (b)
 $T_i \rightarrow T_j \rightarrow T_k \rightarrow T_m$
- (c) – another acceptable serializability order of the transactions: $T_i \rightarrow T_k \rightarrow T_j \rightarrow T_m$.



Examples: Draw Precedence Graphs

- Schedule extra 1:
 - T2:R(A), T1:R(B), T2:W(A), T3:R(A), T1:W(B), T3:W(A), T2:R(B), T2:W(B)



- Schedule extra 2:
 - T2:R(A), T1:R(B), T2:W(A), T2:R(B), T3:R(A), T1:W(B), T3:W(A), T2:W(B)



Examples


□ Schedule extra 1:

- T2:R(A), T1:R(B), T2:W(A), T3:R(A), T1:W(B), T3:W(A), T2:R(B), T2:W(B)



□ Schedule extra 2:

- T2:R(A), T1:R(B), T2:W(A), T2:R(B), T3:R(A), T1:W(B), T3:W(A), T2:W(B)



Data Concurrency and Consistency

- **Data concurrency**

- Many users can access data at the same time.

- **Data consistency**

- Each user sees a **consistent view** of the data
 - Including **visible changes made by the user's own** transactions **and** transactions of **other** users.

4 Types of Concurrency Problems that Lock can Prevent

- Lost updates
 - Two transactions select the same row, and then update the row based on the values originally selected
 - Since each transaction is unaware of the other, the **later update overwrites the earlier** update

4 Types of Concurrency Problems that Lock can Prevent

- Dirty reads
 - A transaction **selects** data that **hasn't been committed** by another transaction
 - If the second transaction rolls back its changes, the first transaction selects a row that does not exist in the database.

4 Types of Concurrency Problems that Lock can Prevent

- Nonrepeatable reads
 - Two **SELECT** statements in **one transaction** that try to get the same data **actually get different values** because **another transaction had updated** the data **during the time between** the two statements

4 Types of Concurrency Problems that Lock can Prevent

- Phantom reads
 - Occurs when **one** transaction **update or delete** a set of rows **at the same time** when **another** transaction does **insert or delete** on the same set of rows
 - Example
 - **TA updates payment_total** for **all invoices** with a balance due
 - **TB inserts new unpaid invoices** in the mean time
 - When TA is done, **invoices with balance due still exist**
 - New unpaid invoices were not updated by TA

The Isolation Models

- Isolation models
 - Controls the degree to which transactions isolate from each other
 - More restrictive level will reduce or eliminate concurrency problems
 - Less restrictive level enhances performance
 - Can prevent different concurrency problems
 - Lost updates
 - Dirty reads
 - Nonrepeatable (fuzzy) reads
 - Phantom reads

Isolation Levels and Concurrency Problems

Isolation level	Problems prevented
READ UNCOMMITTED	None
READ COMMITTED	Dirty reads
REPEATABLE READ	Dirty reads, lost updates, non-repeatable read
SERIALIABLE	All

Oracle isolation levels

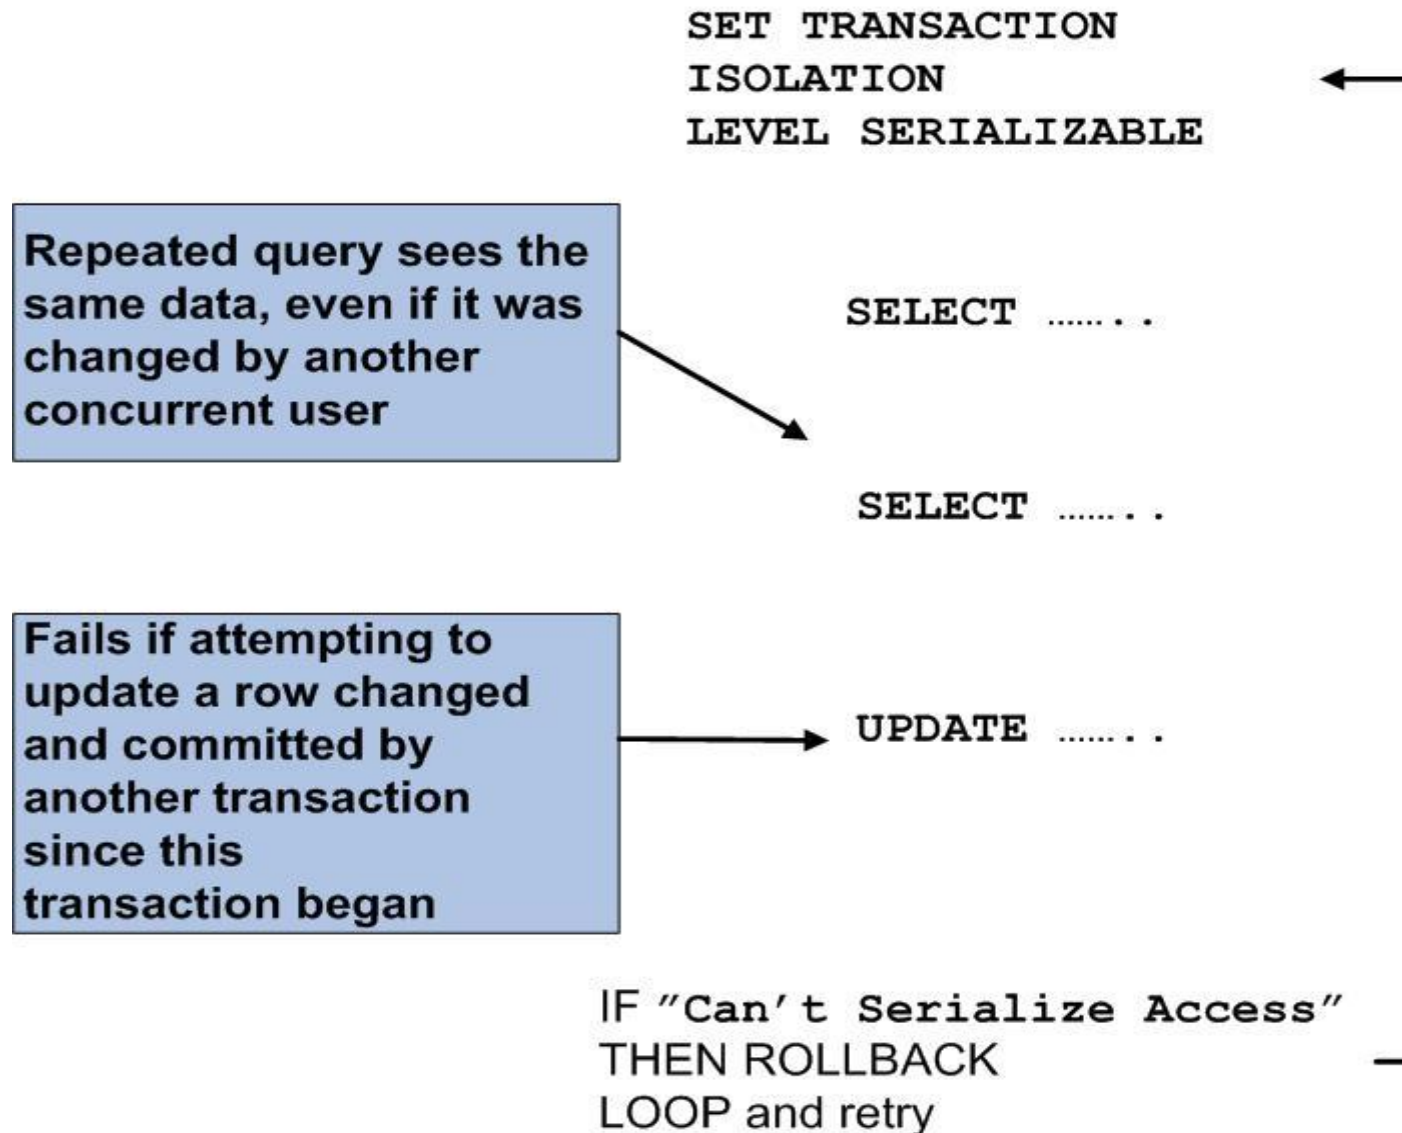
Read committed	Each query executed by a transaction sees only data that was committed before the <u>query</u> began (Oracle default isolation level)
Serializable	Serializable transactions see only those changes that were committed at the time the <u>transaction</u> began , plus its own changes
Read-only	The transaction sees only those changes that were committed at the time the transaction began and do not allow any DML statement

Set the Isolation Level

You can set the isolation level of a transaction by using one of these statements at the beginning of a transaction:

- **SET TRANSACTION ISOLATION LEVEL READ COMMITTED;**
- **SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;**
- **SET TRANSACTION ISOLATION LEVEL READ ONLY;**

Serializable Transaction Failure



Modes of Locking

- **Exclusive lock mode**
 - Prevents the associates resource from being shared.
 - This lock mode is obtained to modify data.
 - The first transaction to lock a resource exclusively is the only transaction that can alter the resource until the exclusive lock is released.

Modes of Locking

- **Share lock mode**
 - **Allows the associated resource to be shared, depending on the operations involved.**
 - **Multiple users reading data can share the data,**
 - **holding share locks to prevent concurrent access by a writer (who needs an exclusive lock).**
 - **Several transactions can acquire share locks on the same resource.**

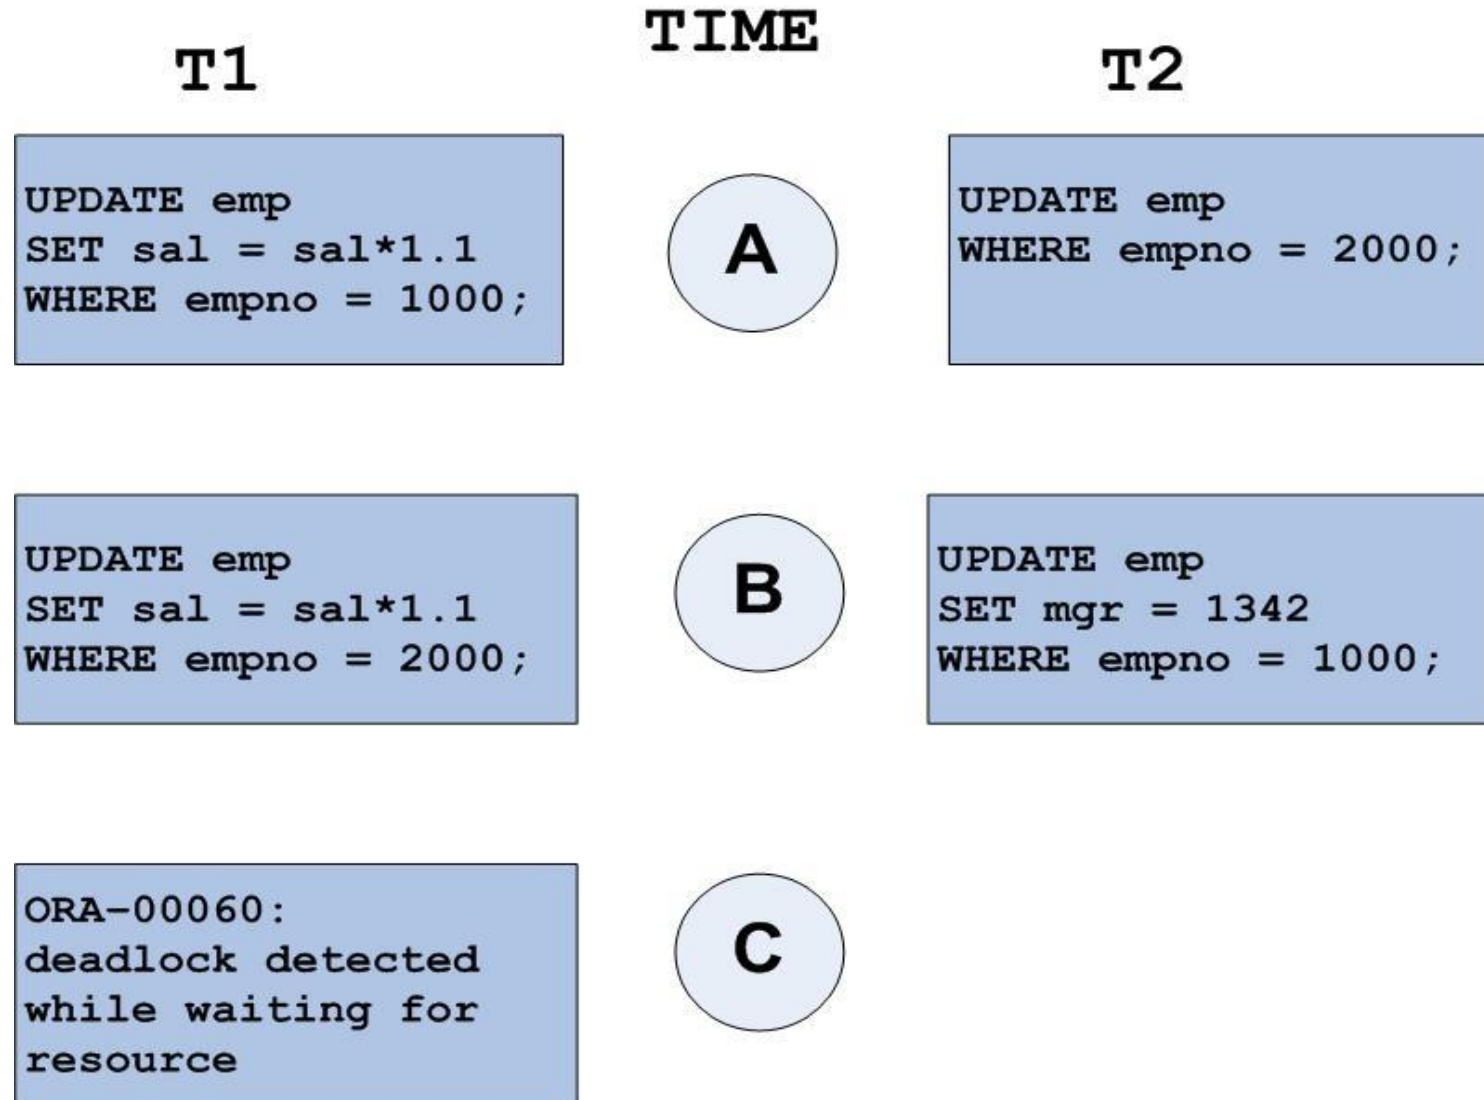
Types of Locks

Lock	Description
DML locks (data locks)	DML locks protect data For example, table locks lock entire tables, rowlocks lock selected rows.
DDL locks (dictionary locks)	DDL locks protect the structure of schema objects
Internal locks and latches	Internal locks and latches protect internal database structures such as datafiles

Deadlock

- **Occurs when neither of two transactions can be committed**
 - **Because each transaction is waiting for the other transaction to release the lock on the resource it needs.**
 - **Example in next slide**

Deadlock



Common Recommendations

- Keep transactions as short as possible
- Avoid executing long-running queries when transactions which update the table are also executing.
- Use lowest possible transaction isolation level
- Make large changes when you are assured of nearly exclusive access